

GALILEO Graph Databases Group

Implementation of GALOIS Features – First Check

Francesco Pivotto 2158296

Giorgia Amato 2159999

Alessio Demo 2142885

December 20, 2025

Contents

1	Iterative Baseline Extension	3
1.1	Motivation and Objective	3
1.2	Design Principles and Architectural Integration	3
1.3	Iterative Execution Workflow	3
1.4	Baseline-Specific Considerations	4
1.5	Outcome and Observations	4
2	Evaluation of the Iterative Baselines	4
2.1	Iterative Execution Results	4
2.2	Comparison with Non-Iterative Baselines	5
3	GALOIS goal and System Architecture	6
3.1	Orchestration and Schema Management	6
3.2	The Execution Engine	6
4	Logical Optimization (Query PLanning)	8
4.1	Predicate Push-Down Strategy	8
4.2	Join Handling Details	8
5	Physical Optimization: Execution Strategies	9
5.1	Table-Scan Algorithm Technical Specification	9
5.1.1	Description	9
5.1.2	Inputs	9
5.2	Key-Scan Algorithm Technical Specification	11
5.2.1	Description	11
5.2.2	Inputs	11
5.3	Adaptive Strategy Selection and Confidence Estimation	13

6	Execution Cycle & Memory Handling	14
6.1	The GaloisMemory Management Class	14
6.1.1	Purpose and Architecture	14
6.1.2	Internal State Variables	14
6.1.3	Core Methods and Functionality	14
6.2	The Iterative Execution Loop and Termination Criteria	15
6.2.1	Loop Mechanism: Pagination & History Injection	15
6.3	Termination Criteria (Convergence Conditions)	16
7	Post-Processing	17
7.1	Robust JSON Repairing	17
7.2	Data Aggregation and Virtualization	17
7.3	Heuristic Type Inference and Sanitization	17
7.4	Deterministic Query Execution	17
8	Numerical Results	19
8.1	GALOIS WO Results comparison	19
8.2	GALOIS S Results comparison	19
8.3	GALOIS A Results comparison	19
8.4	GALOIS F Results comparison	20
9	Experiments	22
9.1	Experiment 1:	22
9.2	Experiment 2: Logical and Physical Optimization Impact	22
9.3	Experiment 3:	23
9.4	Experiment 4:	23
9.5	Experiment 5: Impact of SQL Query Complexity on Retrieval Quality . . .	24
9.6	Experiment 6:	26
9.7	Experiment 7:	26
9.8	Experiment 8:	26
10	Implementation Challenges & Future Work	27
10.1	Performance Deficiencies of the Key-Scan Operator	27
10.1.1	The Locality-of-Reference Problem	27
10.2	Database Resource Management for DuckDB	27
10.3	Addressing Sub-Optimal Prompt Structure	28
11	Contributions	29

1 Iterative Baseline Extension

1.1 Motivation and Objective

Building on the unified execution pipeline introduced in previous phases, this task extends both the SQL and NL baselines with a controlled iterative retrieval mechanism. The goal is to improve result completeness by allowing the LLM to generate multiple partial result sets across successive calls, while preserving full compatibility with the project's standardized JSON output schema.

This mechanism addresses cases where a single LLM invocation fails to recover the entire expected answer set due to truncation, prompt limitations, or generation variability. Iteration provides a systematic way to accumulate non-overlapping tuples across rounds.

1.2 Design Principles and Architectural Integration

The iterative enhancement was introduced without modifying the core architecture. A configurable parameter `max_iter` (defined in `config.yaml`) controls the behavior and ensures consistent execution across providers.

The integration adheres to the following principles:

- **Provider-Agnostic Execution.** Iteration is implemented at the baseline level (SQL and NL), leaving LLM instantiation and LCEL chain construction unchanged. This guarantees full compatibility with any provider supported through the Adapter and Factory patterns.
- **Minimal Invasiveness.** No modifications were made to the LCEL chain, output parser, or prompt template. Iteration is strictly an additional layer wrapped around the existing call-parse-store cycle.
- **Deduplicated Accumulation.** Each iteration contributes new tuples based on strict key-wise comparison. A local accumulator collects all rows returned in previous rounds, preventing redundancy and guaranteeing stable output.
- **Early Stopping Criterion.** To avoid unnecessary LLM calls, the system terminates early if an iteration yields zero new rows. This rule applies identically to SQL and NL baselines.

1.3 Iterative Execution Workflow

For each query (SQL baseline) or natural language prompt (NL baseline), the extended workflow proceeds as follows:

1. **Initialization.** Create an accumulator for unique tuples and load `max_iter`.
2. **Repeated Invocation.** At each iteration, the baseline executes:
 - (a) construction of the LCEL chain with the existing prompt template,
 - (b) LLM invocation using the configured provider,
 - (c) parse the response via the Pydantic output parser.

3. **Deduplication and Union.** Compare new tuples to the accumulator using canonical JSON formatting; retain only unseen tuples.
4. **Stopping Condition.** Terminate when no new tuples are found.
5. **Persistence.** After early stopping or reaching `max_iter`, serialize:
 - the union of all retrieved tuples,
 - total elapsed time,
 - aggregated token usage.

1.4 Baseline-Specific Considerations

SQL Baseline. Iteration compensates for incomplete tuple enumeration in a single pass. Because SQL inputs are fully specified, repeated calls primarily improve recall.

NL Baseline. Iteration mitigates incomplete enumerations and semantic drift in free-form queries. Early stopping is essential to prevent hallucinated or reformulated rows once the LLM stabilizes on a fixed subset.

1.5 Outcome and Observations

Controlled iteration increases tuple coverage while remaining fully aligned with the evaluation framework. The design preserves modularity, provider-agnosticism, and the standardized JSON format established in previous phases.

The mechanism also maintains conceptual continuity with iterative operators in the GALOIS framework (e.g., Key-Scan), providing a foundation for more advanced retrieval strategies in future work.

2 Evaluation of the Iterative Baselines

We evaluated the iterative SQL and NL baselines on the full benchmark suite (FLIGHT-2, FLIGHT-4, GEO, MOVIES, PRESIDENTS, WORLD).

2.1 Iterative Execution Results

First of all for what concerns the `llama-3.3-70b-instruct` model, the following table highlights the differences in terms of metrics between results presented in the original GALOIS paper and our implementation with iterative extension:

Maintaining consistency with the previous report, the following table show the results obtained with the iterative extension for both baselines using `gemini-2.5-flash` models:

The iterative execution impacts the behavior of both baselines in distinct ways. In the SQL setting, the model benefits from task repetition, often recovering additional valid tuples that were missed in the first pass. This results in higher tuple-level completeness and, in many cases, improvements in the F1-CELL and CARDINALITY metrics. For NL queries, which are inherently more open-ended, the iterative mechanism still enhances tuple coverage, though with a less uniform effect across datasets.

Metric	PAPER	Ours	$\Delta(\text{PAPER} - \text{Ours})$
F1-Cell	0.431	0.403	0.028
Cardinality	0.659	0.465	0.194
Tuple Constr.	0.351	0.197	0.154
AVG-Score	0.481	0.355	0.126
#Tokens (M)	0.330	0.008	0.322
Avg Time	61.400	5.263	56.137

Table 1: sql_baseline

Metric	PAPER	Ours	$\Delta(\text{PAPER} - \text{Ours})$
F1-Cell	0.237	0.271	-0.034
Cardinality	0.462	0.500	-0.038
Tuple Constr.	0.065	0.091	-0.026
AVG-Score	0.254	0.287	-0.033
#Tokens (M)	0.830	0.005	0.825
Avg Time	120.000	2.206	117.794

Table 2: nl_baseline

Metric	NL	SQL
F1-CELL	0.425	0.456
CARDINALITY	0.696	0.607
TUPLE-CONSTR.	0.262	0.245
#TOKENS (M)	0.028	0.067
AVG-TIME (s)	5.035	9.681
AVG-SCORE	0.461	0.436

Table 3: Iterative Results for gemini-2.5-flash.

2.2 Comparison with Non-Iterative Baselines

To contextualize the effect of iteration, we compare the results above with the single-pass baselines. Overall, the iterative variant provides measurable gains in tuple completeness and global AVG-SCORE, especially in datasets where the model tended to omit large fractions of the answer in a single call. As expected, the iterative mechanism also increases the total number of generated tokens and the average execution time. However, these costs remain moderate and proportional to the configured value of `max_iter`.

In conclusion, iterative execution improves robustness and recall, providing a more stable foundation for evaluating LLM behavior under both SQL and NL interfaces.

3 GALOIS goal and System Architecture

Large Language Models (LLMs) can answer declarative queries expressed in Natural Language (NL) or SQL, but their one-shot responses frequently fail to recover complete and factually consistent tuples when queries involve filtering, joins, or aggregations. As highlighted in the GALOIS paper, direct prompting exhibits limited recall and incomplete enumeration, making LLMs unreliable as standalone query processors.

The GALOIS framework mitigates these limitations by treating the LLM as a latent knowledge store and decomposing execution into logical and physical operators, in analogy with traditional DBMSs. Rather than issuing a single prompt, GALOIS executes structured scan and filtering strategies (Table Scan, Key Scan) together with multiple logical planning variants (push-all, push-selective, push-wo, and the optimized GALOIS-F). This operator-based design improves tuple completeness and stabilizes execution across heterogeneous query types.

Our implementation orchestrates the interaction between a declarative SQL interface and the probabilistic nature of LLMs through four distinct stages: parsing, logical planning, physical execution, and local post-processing. This section details the core components implemented in the `Galois` orchestration class.

3.1 Orchestration and Schema Management

The entry point of the system is the `Galois` class, which initializes the workflow by parsing the input SQL query into an Abstract Syntax Tree (AST) using the `sql_query_parser` module. Unlike standard text-to-SQL approaches, GALOIS does not immediately prompt the LLM with the raw query, instead it inspects the database structure using the `GaloisSchemaManager`. This component connects to a read-only DuckDB instance to retrieve accurate metadata (column names, primary keys, and data types). This schema awareness allows the system to construct semantically valid JSON schemas for the prompts, ensuring that the LLM’s output adheres to the strict typing required for the subsequent relational operations.

Data-Driven Schema Enrichment: The Schema Manager has been enhanced to support few-shot context injection. For textual columns (VARCHAR/TEXT) that are not primary keys, the manager queries the underlying DuckDB instance to retrieve a small sample of actual values (e.g., 3 examples). These examples are injected directly into the JSON schema description passed to the LLM ("description": "Examples: ..."), providing the model with immediate context about data formatting and distribution.

3.2 The Execution Engine

The retrieval logic is encapsulated within the `GaloisExecutor` class. This component abstracts the complexity of the LLM interaction (via `LangChain`) and implements the iterative retrieval loops.

State Management and Deduplication: To handle the context window limitations of LLMs, the executor utilizes a custom `GaloisMemory` class. This memory component maintains a history of retrieved tuples and implements a strict deduplication mechanism based on primary key hashing (`_get_key_values`). This ensures that subsequent iterations (up to a configurable `max_iter`) prompt the model to generate only novel data.

Iterative Scan Operators: The executor supports two physical operators:

- **Table-Scan:** Iteratively requests complete tuples (keys and attributes) until convergence or iteration limits are reached.
- **Key-Scan:** Decomposes retrieval into two phases. First, it iterates to collect all unique primary keys. Once the keys are consolidated, it issues a final "Tuple Request" to retrieve non-key attributes for the collected identifiers.

4 Logical Optimization (Query PLanning)

The Logical Optimization phase determines what information needs to be extracted from the LLM and which filtering conditions should be "pushed down" to the generative model versus applied locally. This decision-making process is critical to balancing the trade-off between retrieval precision and the risk of hallucination.

4.1 Predicate Push-Down Strategy

The system implements a flexible planning mechanism capable of generating different execution plans based on the configured optimization level. The planner analyzes the **WHERE** clause of the parsed SQL query and partitions the conditions into two sets: **conditions_to_push** (integrated into the prompt) and **residual_conditions** (deferred to the post-processor). The framework supports four distinct logical variants:

- **No-Push (GALOIS-WO)**: Treats the LLM purely as a data source. No filtering conditions are included in the prompt; the system retrieves the entire extension of the table (or its keys) using the key-scan and performs all filtering locally.
- **Push-All (GALOIS-A)**: Maximizes delegation by translating all SQL predicates into natural language constraints within the prompt. This assumes high model capability in understanding and applying complex logic.
- **Push-Selective (GALOIS-S)**: A heuristic approach that identifies and pushes only the most selective condition, using the LLM's confidence self-assessment as a proxy to determine which conditions are "safe" or "selective" enough to be handled directly by the generative model and aims to reduce the search space without overwhelming the model context.
- **Push-Confident (GALOIS-F)**: An adaptive strategy that utilizes a **ConfidenceEstimator**. The system queries the LLM to self-evaluate its confidence in correctly assessing specific predicates (returning a "HIGH" or "LOW" score). Only conditions with high confidence scores are pushed down, while ambiguous ones are retained for local execution. Furthermore the LLM also does an estimation process over the Physical optimization to apply looking to the entire query, the details will be discussed in section 5.3.

4.2 Join Handling Details

For multi-table queries, the logical planner decomposes the join operation into independent scan tasks. The system identifies all tables involved in the query (FROM clause and JOIN clauses) and generates a sub-plan for each. The planner ensures that alias references in the SQL query are correctly mapped to their respective table scans, allowing the Executor to fetch data for each relation independently before they are rejoined in the Post-Processing phase.

5 Physical Optimization: Execution Strategies

While Logical Optimization defines the declarative scope of the data retrieval task, Physical Optimization is responsible for specifying the concrete, algorithmic execution strategy. The proposed system implements a set of iterative data extraction algorithms designed to mitigate the inherent context window limitations of Large Language Models (LLMs) and ensure comprehensive, de-duplicated data acquisition.

5.1 Table-Scan Algorithm Technical Specification

5.1.1 Description

The **Table-Scan** algorithm is a single-phase, iterative extraction technique that leverages an LLM to retrieve a complete collection of unique, full-schema records (tuples) satisfying the input Q . This operator is designed to maximize coverage by focusing on the generation and accumulation of entire records in each iteration step, contrasting with multi-phase strategies.

1. **Iterative Tuple Collection** (Algorithm 1, Lines 4–19): The algorithm executes a loop, repeatedly invoking the LLM chain to discover and accumulate novel, complete unique tuples that conform to the constraints of Q .
 - The initial invocation utilizes a base input prompt derived from Q .
 - Subsequent iterations, $i > 0$, dynamically modify the input request by injecting the historical context ($\mathcal{H}$) of previously accumulated tuples. This mechanism guides the LLM to seek novel, uncollected records, thereby enforcing deduplication and maximizing the result set.
 - The process is bounded by I_{max} and features a termination condition: if the post-response parsing step confirms the absence of new unique tuples ($\mathcal{R}_{new} = \emptyset$), the loop is prematurely exited (handling the Convergence criterion).
2. **Finalization and Cleanup** (Lines 22–24): Upon termination, the consolidated set of unique tuples residing in the internal $\mathcal{M}$ (Memory) structure is designated as the $\mathcal{R}_{final}$. The state of the internal memory/history is subsequently reset.

5.1.2 Inputs

- Q : The declarative data retrieval statement (e.g., an SQL-like query).
- $\mathcal{P}_{pushdown}$: An optional set of **WHERE** clause predicates pushed down to refine the extraction scope.
- I_{max} : An optional integer parameter specifying the maximum number of LLM invocations.

Algorithm 1 Table-Scan Algorithm

```
1: function TABLE_SCAN( $Q, \mathcal{P}_{pushdown}, I_{max}$ )
2:    $\mathcal{I} \leftarrow \text{BUILDINITIALREQUEST}(Q)$ 
3:    $\mathcal{C} \leftarrow \text{INITIALIZELLMINTERACTIONCHAIN}$ 
4:    $i \leftarrow 0$ 
5:    $\mathcal{H} \leftarrow \emptyset$ 
6:
7:   while  $i < I_{max}$  do
8:     if  $i = 0$  then
9:        $\mathcal{R}_{raw} \leftarrow \mathcal{C}.\text{INVOKE}(\mathcal{I})$ 
10:    else
11:       $\mathcal{I} \leftarrow \text{UPDATEREQUESTWITHHISTORY}(\mathcal{H})$ 
12:       $\mathcal{R}_{raw} \leftarrow \mathcal{C}.\text{INVOKE}(\mathcal{I})$ 
13:
14:       $\mathcal{R}_{parsed} \leftarrow \text{PARSELMLRESPONSE}(\mathcal{R}_{raw})$ 
15:       $\mathcal{R}_{new} \leftarrow \text{IDENTIFYNEWUNIQUETUPLES}(\mathcal{R}_{parsed}, \mathcal{H})$ 
16:      if  $\mathcal{R}_{new} = \emptyset$  then
17:        break
18:
19:       $\text{ADDTOMEMORYANDHISTORY}(\mathcal{R}_{new}, \mathcal{M}, \mathcal{H})$ 
20:       $i \leftarrow i + 1$ 
21:
22:   $\mathcal{R}_{final} \leftarrow \text{RETRIEVEALLTUPLES}(\mathcal{M})$ 
23:   $\text{CLEARSTATE}(\mathcal{M}, \mathcal{H})$ 
24:
25:  return  $\mathcal{R}_{final}$ 
```

5.2 Key-Scan Algorithm Technical Specification

5.2.1 Description

The Key-Scan algorithm constitutes a specialized, two-phase iterative data extraction strategy. It employs an LLM to retrieve the required record set by prioritizing the discovery and consolidation of unique primary/candidate key values before performing a final request for associated, non-key attributes. This two-phase approach inherently ensures record uniqueness and may improve token efficiency by minimizing the size of the history context injected into the prompt.

1. Iterative Key Collection (Algorithm 2, Lines 4–19): This phase iteratively invokes the LLM chain to discover and accumulate a complete, unique set of $\mathcal{K}$ (key values) relevant to Q .
 - The first iteration uses a prompt formulated to extract only the key attributes.
 - Subsequent iterations modify the input request by embedding a history ($\mathcal{H}$) of previously collected key values, steering the LLM’s generation toward novel, uncollected records.
 - The loop is governed by I_{max} and terminates if the parsing step confirms that no new unique key values ($\mathcal{K}_{new} = \emptyset$) were returned by the LLM (Convergence).
2. Final Attribute Retrieval (Tail Request) (Lines 22–26): Upon key collection convergence, the consolidated set of $\mathcal{K}_{all}$ is compiled. A single, dedicated final request (the *Tail Request*) is issued to the LLM to efficiently fetch all missing non-key attributes ($\mathcal{A}$) corresponding to the collected keys.

The function concludes by performing a relational join operation between the collected keys and the attributes from the final response to construct the $\mathcal{R}_{final}$ of complete, de-duplicated records.

5.2.2 Inputs

- Q : The SQL-like data retrieval statement.
- $\mathcal{P}_{pushdown}$: A list of WHERE clause predicates for extraction refinement.
- I_{max} : An optional integer parameter bounding the iterative key collection phase.

Algorithm 2 Key-Scan Algorithm

```
1: function KEY_SCAN( $Q, \mathcal{P}_{pushdown}, I_{max}$ )
2:    $\mathcal{I} \leftarrow \text{BUILDINITIALREQUEST}(Q)$ 
3:    $\mathcal{C} \leftarrow \text{INITIALIZELLMINTERACTIONCHAIN}$ 
4:    $i \leftarrow 0$ 
5:    $\mathcal{H} \leftarrow \emptyset$ 
6:
7:   while  $i < I_{max}$  do
8:     if  $i = 0$  then
9:        $\mathcal{R}_{raw} \leftarrow \mathcal{C}.\text{INVOKE}(\mathcal{I})$ 
10:    else
11:       $\mathcal{I} \leftarrow \text{UPDATEREQUESTWITHHISTORY}(\mathcal{H})$ 
12:       $\mathcal{R}_{raw} \leftarrow \mathcal{C}.\text{INVOKE}(\mathcal{I})$ 
13:
14:       $\mathcal{K}_{parsed} \leftarrow \text{PARSELLMRESPONSE}(\mathcal{R}_{raw})$ 
15:       $\mathcal{K}_{new} \leftarrow \text{IDENTIFYNEWUNIQUEKEYS}(\mathcal{K}_{parsed}, \mathcal{H})$ 
16:      if  $\mathcal{K}_{new} = \emptyset$  then
17:        break
18:
19:       $\text{SAVEKEYVALUESANDUPDATEHISTORY}(\mathcal{K}_{new}, \mathcal{H})$ 
20:       $i \leftarrow i + 1$ 
21:
22:    $\mathcal{K}_{all} \leftarrow \text{RETRIEVEALLCOLLECTEDUNIQUEKEYS}(\mathcal{H})$ 
23:    $\mathcal{I}_{final} \leftarrow \text{BUILDFINALATTRIBUTEREQUEST}(\mathcal{K}_{all})$ 
24:    $\mathcal{R}_{final\_raw} \leftarrow \mathcal{C}.\text{INVOKE}(\mathcal{I}_{final})$ 
25:    $\mathcal{A}_{parsed} \leftarrow \text{PARSEFINALRESPONSE}(\mathcal{R}_{final\_raw})$ 
26:    $\mathcal{R}_{final} \leftarrow \text{RELATIONALJOIN}(\mathcal{K}_{all}, \mathcal{A}_{parsed})$ 
27:    $\text{CLEARSTATE}(\mathcal{H})$ 
28:
29:   return  $\mathcal{R}_{final}$ 
```

5.3 Adaptive Strategy Selection and Confidence Estimation

The framework incorporates an adaptive execution layer that dynamically mediates the selection between the Table-Scan and Key-Scan operators. This decision is parameterized by a Confidence Estimator, which quantitatively assesses the expected factual data retrieval reliability based on the Q 's schema complexity and the LLM's raw confidence score.

The confidence estimation model incorporates a decay factor sensitive to the number of projected columns. The final confidence score C_{final} is calculated from the LLM's intrinsic raw confidence (C_{raw}) by applying an exponential decay based on N_{cols} , the cardinality of attributes in the SELECT clause:

$$C_{final} = (C_{raw})^{N_{cols}}$$

This formulation rigorously models the non-linear increase in difficulty for the LLM to generate a coherent, fully populated tuple set as the schema width grows.

Execution Strategy Selection Heuristic:

- If C_{final} exceeds a predefined operational threshold (τ), the system prioritizes the Key-Scan approach. This two-phase method aligns with a Chain-of-Thought reasoning process, which is generally associated with higher accuracy and factuality when the model exhibits sufficient confidence.
- Conversely, if C_{final} falls below τ , or if the schema is judged to be excessively wide ($N_{cols} \gg 1$), the framework defaults to the Table-Scan operator. Table-Scan is considered the more robust strategy when the factual confidence is low, as it reduces the complexity of the initial extraction target (the full tuple vs. just the key).

This adaptive mechanism is fundamental for optimizing the execution of the full GALOIS-F logical plan by selecting the most appropriate physical operator based on a data-driven confidence metric.

6 Execution Cycle & Memory Handling

The interaction between the deterministic orchestration layer and the generative model is governed by a stateful execution cycle. Unlike standard single-turn LLM interactions, the Galois framework maintains a persistence layer across iterations to maximize recall and prevent redundancy. This logic is encapsulated within the **GaloisExecutor** class and its specialized internal memory component, **GaloisMemory**.

6.1 The GaloisMemory Management Class

6.1.1 Purpose and Architecture

The **GaloisMemory** class serves as the dedicated state management component for the physical execution strategies (**Table-Scan** and **Key-Scan**). Its primary responsibilities are:

1. **History Accumulation:** Storing the unique records (tuples) extracted by the LLM chain across multiple iterative steps.
2. **Deduplication Enforcement:** Ensuring that only unique records are preserved in the memory structure ($\mathcal{M}$) to maintain data integrity and prevent redundant processing.
3. **Context Provision:** Formatting the accumulated records ($\mathcal{M}$) into a structured history ($\mathcal{H}$) for injection into subsequent LLM prompts, guiding the model toward novel data discovery.

The class inherits from **BaseMemory** and **BaseModel**, structuring the internal state via Pydantic fields.

6.1.2 Internal State Variables

The memory state is maintained by the following core attributes:

- **memory** (`List[Dict[str, Any]]`): The primary storage structure containing the complete collection of unique records (tuples) retrieved thus far.
- **key_values** (`set`): A hash set storing the unique key value combinations (as Python tuples) for rapid, $O(1)$ deduplication lookups. This structure is central to enforcing uniqueness across iterations.
- **tokens** (`int`): An aggregate counter tracking the total number of tokens processed by the LLM across all invocations for the current extraction task.
- **time** (`float`): An aggregate measure of the total processing time (latency) associated with the LLM interactions.

6.1.3 Core Methods and Functionality

1. **load_memory_variables(inputs)**
 - *Function:* Prepares the accumulated records for use as historical context in the next LLM prompt.

- *Implementation:* Converts the internal `memory` list into a compact JSON string, which is returned under the key `"history"`. This serialized format is robust for prompt injection.

2. `save_context(inputs, outputs)`

- *Function:* The central method for processing new records generated by an LLM invocation and updating the memory state.
- *Procedure:* Iterates through the `outputs["response"]` (newly generated records). For each record, it extracts the unique key values using `_get_key_values`. If the key is not present in `self.key_values`, the record is appended to `self.memory`, and the key is added to `self.key_values`.
- *Convergence Check:* Compares the size of `self.key_values` before and after processing the new records. If the set size remains unchanged, it indicates that *no new unique tuples* were found, triggering a `NoNewTuplesFound` exception to terminate the iterative process (per Algorithms 1 and 2).
- *Metric Update:* Aggregates the `tokens` and `time` metrics from the `outputs` into the class state.

3. `clear()`

- *Function:* Resets the memory state, typically called after an extraction process completes.
- *Implementation:* Empties the `self.memory` list and `self.key_values` set, re-setting `tokens` and `time` to zero.

4. `_get_key_values(record, keys)`

- *Function:* An internal helper method to extract the composite key values from a single record dictionary.
- *Implementation:* Takes a `record` (dict) and a list of `keys` (column names) and returns a `tuple` containing the corresponding values. This tuple is used as the hashable item for deduplication checks in the `key_values` set.

6.2 The Iterative Execution Loop and Termination Criteria

Both the **Table-Scan** and **Key-Scan** physical operators rely on a fundamental Iterative Execution Loop (lines 4-19 in Algorithms 1 and 2) to overcome the LLM context window limitations and maximize data extraction coverage. The core mechanism involves repeatedly querying the LLM and dynamically updating the prompt with historical data.

6.2.1 Loop Mechanism: Pagination & History Injection

In both strategies, the iterative process is driven by the dynamic insertion of previously acquired data into the subsequent prompt. This mechanism ensures two critical functions:

1. **Deduplication:** The LLM is implicitly guided to avoid regenerating records already present in the injected history ($\mathcal{H}$).

2. **Novelty Search:** The context directs the LLM to explore uncollected portions of the solution space, thereby maximizing the cardinality of the final result set ($\mathcal{R}_{final}$).
3. To enhance retrieval efficiency and prevent the LLM from wandering randomly, the iterative mechanism implements a **value-based pagination** strategy: In addition to the history the system identifies the last retrieved value (`last_val`) from the memory. The prompt logic explicitly instructs the LLM to generate unique tuples that strictly follow `last_val` (alphabetically or numerically). This transforms the extraction into a pseudo-sequential scan, improving coverage compared to simple history injection.

The difference lies in the content of the history: **Table-Scan** injects a history of complete tuples, while **Key-Scan** injects only the unique key values ($\mathcal{K}$) collected so far.

6.3 Termination Criteria (Convergence Conditions)

The iterative loop is governed by two distinct termination criteria, ensuring both safety against infinite execution and functional completeness (convergence). The loop terminates if either condition is met:

1. Maximum Iteration Limit (I_{max}):

- *Condition:* The loop counter (i) reaches the maximum permissible value defined by the input parameter I_{max} , which is configurable via the system’s YAML configuration file or directly supplied by the user.
- *Purpose:* This serves as a fail-safe mechanism to bound the token consumption and execution time, preventing excessive resource usage in cases where the LLM continually produces marginally unique or redundant data.

2. Data Convergence ($\mathcal{R}_{new} = \emptyset$ or $\mathcal{K}_{new} = \emptyset$):

- *Condition:* After an LLM invocation and subsequent parsing/deduplication step, the system determines that **no new unique records (tuples)** were added to the memory (for **Table-Scan**), or **no new unique key values** were collected (for **Key-Scan**).
- *Mechanism:* This condition is detected by the `save_context` function of the `GaloisMemory` class, typically by comparing the cardinality of the `key_values` set before and after processing the LLM’s latest output.
- *Purpose:* This is the primary functional exit criterion, indicating that the LLM has likely exhausted its capacity to discover novel data relevant to the query Q , signaling functional convergence and allowing for premature loop termination for efficiency.

The system employs the `break` statement within the `While` loop when convergence is detected, overriding the I_{max} constraint if full coverage is achieved earlier.

7 Post-Processing

The Post-Processing module serves as the deterministic convergence point of the architecture, addressing the inherent impedance mismatch between the probabilistic output of the Large Language Model and the strict requirements of the SQL standard. This phase transforms the raw tuple sets retrieved by the Executor into a consistent relational state, enabling the execution of complex logical operations that cannot be reliably offloaded to the generative model.

7.1 Robust JSON Repairing

To address the issue of invalid JSON syntax produced by LLMs (e.g., trailing commas, missing null values, or markdown code blocks), we implemented a deterministic sanitization routine (`repair_json_content`) within the Executor. Before parsing, the raw response creates a sanitized string using regular expressions to strip Markdown tags and fix common syntax errors (such as converting empty assignments `key: ,` to `key: null,`). This preprocessing step significantly reduced parsing failures.

7.2 Data Aggregation and Virtualization

Upon the completion of the physical scan phase (whether Table-Scan or Key-Scan), the retrieved tuples are aggregated into a local staging structure, referred to as the **"Data Lake"** which maps table names to lists of dictionaries. To bridge the gap to a relational environment, the system instantiates an ephemeral, in-memory DuckDB connection. The aggregated data is converted into Pandas DataFrames and registered as virtual views within this database instance, effectively creating a transient relational schema populated by the LLM-generated data.

7.3 Heuristic Type Inference and Sanitization

The system implements a strict type inference policy to preserve data integrity. The `perform_local_join_and_query` method employs specific regular expressions to identify and normalize various string representations of null produced by LLMs (e.g., `'nan'`, `'null'`, `'n/a'`) into native Python `None` objects. Furthermore, a strict casting policy is applied via Pandas: columns are converted to numeric types only if the conversion preserves data integrity (checking the ratio of numeric vs. non-null values); otherwise, they revert to object types to avoid data loss during the virtualization in DuckDB. This zero-tolerance approach ensures that alphanumeric identifiers or mixed-format codes are not erroneously coerced into null values.

7.4 Deterministic Query Execution

Once the data is virtualized and typed, the system executes the logical query plan locally. This step involves running the original SQL query (sanitized of specific schema prefixes like `.target`) against the virtual tables. This local execution performs two vital functions:

- **Join Reconstruction:** It re-establishes the relational links between disparate datasets that were retrieved via independent scan operations, ensuring the correct implementation of INNER or LEFT joins defined in the original AST.

- **Residual Filtering:** It applies the **residual_conditions** predicates identified during the Logical Optimization phase as too complex or numerically precise for the LLM (e.g., strict inequalities like $\text{length_in_km} > 750$). This ensures that the final result set adheres strictly to the logical constraints of the user’s query, filtering out any hallucinations or imprecise data artifacts introduced by the model.

8 Numerical Results

In this section we will show a comparison between the metrics of original GALOIS paper written by the authors, and the metrics that we obtained from our GALOIS implementation.

8.1 GALOIS WO Results comparison

Metric	PAPER	Ours	$\Delta(\text{PAPER} - \text{Ours})$
F1-Cell	0.518	0.157	0.361
Cardinality	0.691	0.357	0.334
Tuple Constr.	0.389	0.060	0.329
AVG-Score	0.531	0.191	0.340
#Tokens (M)	19.710	0.163	19.709
Avg Time	1460.0	43.103	1416.897

Table 4: GaloisWO comparison

8.2 GALOIS S Results comparison

Metric	PAPER	Ours	$\Delta(\text{PAPER} - \text{Ours})$
F1-Cell	0.480	0.434	0.046
Cardinality	0.655	0.667	-0.012
Tuple Constr.	0.365	0.297	0.068
AVG-Score	0.500	0.466	0.034
#Tokens (M)	0.960	0.088	0.872
Avg Time	130.00	17.952	112.048

Table 5: GaloisS comparison

8.3 GALOIS A Results comparison

Metric	PAPER	Ours	$\Delta(\text{PAPER} - \text{Ours})$
F1-Cell	0.543	0.417	0.126
Cardinality	0.799	0.614	0.185
Tuple Constr.	0.448	0.279	0.169
AVG-Score	0.592	0.436	0.156
#Tokens (M)	0.950	0.076	0.874
Avg Time	120.50	24.481	96.019

Table 6: GaloisA comparison

8.4 GALOIS F Results comparison

Metric	PAPER	Ours	$\Delta(\text{PAPER} - \text{Ours})$
F1-Cell	0.563	0.319	0.244
Cardinality	0.835	0.528	0.307
Tuple Constr.	0.464	0.166	0.298
AVG-Score	0.622	0.338	0.284
#Tokens (M)	1.720	0.142	1.578
Avg Time	47.400	26.196	21.204

Table 7: GaloisF comparison

As we can see from the tables above, we obtained the worst results from GALOIS WO version due to three factors: the probabilistic nature of "Context-Free" retrieval, the loss of semantic cohesion in the Key-Scan operator, and the depth of iteration in simulated table scans.

- Context-Free Retrieval and Popularity Bias:** The primary cause of the low recall in GALOIS WO is the absence of predicate pushdown. By design, GALOIS WO forces the logical plan to strip all **WHERE** clauses, presenting the LLM with a generic query (e.g., "List all presidents" instead of "List presidents of Venezuela"). Unlike a deterministic database storage engine that scans all records faithfully, an LLM is a probabilistic engine heavily influenced by popularity bias. Without the "guidance" of specific filters in the prompt, the model generates general entities. Since the GaloisPostProcessor can only filter data that has been successfully retrieved analyzing the former query, the lack of relevant entries renders the local filtering phase ineffective.
- Semantic Decohesion in the Key-Scan Operator:** Our implementation highlights a structural weakness in the Key-Scan physical operator, in details it decouples the retrieval of keys (Phase 1) from their attributes (Phase 2). In a standard Table-Scan (used in GALOIS A/S), the co-occurrence of related tokens in the generation window (e.g., generating "Venezuela" immediately next to a candidate name) helps the LLM's self-attention mechanism maintain semantic consistency. By isolating the key generation, the Key-Scan loses this context, furthermore in a "No-Push" scenario, the first phase wastes significant token budget generating irrelevant keys, which are then discarded during post-processing, drastically reducing the system's overall efficiency.
- The Trade-off Between Execution Time and Scan Depth:** A critical divergence from the reference paper is observed in the average execution time. This discrepancy indicates that our implementation imposes a stricter limit on iterations (`max_iter` defaults to 4 in `executor.py`). To successfully simulate a "Full Table Scan" on an LLM—essentially brute-forcing the model to output its entire knowledge base for a given entity—a system must iterate significantly longer to bypass the initial layer of popular results. Our lower execution time suggests that the retrieval process is terminated prematurely, capturing only the "surface" data. This suggests that GALOIS WO is theoretically viable only with massive iteration depth, which may be inefficient for real-world applications due to latency and token cost

constraints. Conversely, GALOIS S and A perform better because the prompt acts as a "soft index" allowing the LLM to jump directly to the relevant data partition without requiring exhaustive iteration.

9 Experiments

9.1 Experiment 1:

9.2 Experiment 2: Logical and Physical Optimization Impact

This experiment evaluates the effectiveness of the logical and physical optimization strategies adopted by GALOIS, following the methodology described in Sections 4 and 5 of the original paper. The goal is to isolate and quantify how different execution strategies affect result quality, measured through the **AVG-Score** (average of F1-Cell, Cardinality, and Tuple Constraint).

All experiments are conducted on the **GEO** dataset, which queries the internal parametric knowledge of the LLM and includes 32 SQL queries with pre-computed ground truth.

Physical Optimization. At the physical level, GALOIS selects how data is retrieved from the LLM. For each query, we compare three execution strategies:

- **Table-Scan:** retrieves full tuples in a single scan, providing broader context to the LLM.
- **Key-Scan:** first retrieves key values and then populates remaining attributes, using a chain-of-thought style interaction.
- **AUTO:** lets GALOIS automatically choose between Table-Scan and Key-Scan based on the confidence estimator described in the paper.

We report both the average result quality and the total token consumption. Additionally, we evaluate the *AUTO accuracy*, i.e., how often the automatically selected strategy matches the best-performing fixed strategy (Table-Scan or Key-Scan) for a given query.

Table 8: Physical optimization results on GEO (AVG-Score and token usage).

Strategy	AVG-Score	Tokens (M)
Table-Scan	0.315	0.030
Key-Scan	0.272	0.033
AUTO	0.274	0.031

The results show that, on GEO, **Table-Scan achieves the highest average quality**, suggesting that richer contextual prompts help the LLM retrieve more complete and accurate data. The AUTO strategy closely follows the best fixed plan, confirming that the confidence-based physical planner behaves robustly, although the gain over Table-Scan is limited for this dataset.

Logical Optimization. To isolate the impact of logical planning decisions, we fix the physical strategy to **Table-Scan** and evaluate different predicate pushdown strategies, as described in the paper:

- **NO-PUSH**: no conditions are pushed down into the LLMScan.
- **PUSH-ALL** (GALOIS_A): all WHERE conditions are pushed down.
- **PUSH-CONFIDENT** (GALOIS_F): only conditions selected by the confidence estimator are pushed down.

Only queries with at least two WHERE conditions are considered, in accordance with the experimental protocol of the original work.

Table 9: Logical optimization results on GEO with fixed Table-Scan.

Strategy	AVG-Score	Tokens (M)
NO-PUSH	0.318	0.015
GALOIS A (push all)	0.274	0.009
GALOIS F (confident)	0.397	0.014

The results clearly show that **confidence-based pushdown** (GALOIS_F) significantly outperforms both NO-PUSH and PUSH-ALL strategies. This confirms the key intuition of the paper: pushing all predicates may reduce cost but often degrades result quality due to increased reasoning complexity for the LLM, while selectively pushing only confident predicates yields a better balance between precision and recall.

Comparison with the Paper. Our findings are fully consistent with the original GALOIS results. In particular, we observe the same qualitative trends:

- (i) Table-Scan can outperform Key-Scan when additional context improves factual retrieval, and
- (ii) confidence-driven logical optimization (GALOIS_F) consistently achieves the best AVG-Score among logical strategies.

These results further validate the robustness of the GALOIS optimization framework across datasets and implementations.

9.3 Experiment 3:

9.4 Experiment 4:

Here we implemented an automated evaluation pipeline ("Impact of Rarity and Temporality"), focusing on the **PRESIDENTS** domain. The primary objective is to quantify how the frequency of data in pre-training (popularity) and its temporal recency influence the ability of Large Language Models (LLMs) to measure two specific biases.

- **Rarity**: Compares performance on popular entities (USA) versus rare entities (Venezuela).
- **Temporality**: Compares performance on recent events versus historical events. For this purpose we split the queries into three categories: ALL-TIME covering all dates, RECENT for queries that cover data from late 1900 to today, and PAST for queries for data till the late 1900.

By benchmarking model modalities (GALOIS, NL, SQL) against the Ground Truth, the script calculates the AVG-SCORE. The results confirm that query generation quality degrades significantly when dealing with rare or historical data, as we can see in the tables below:

Table 10: Experimental Results: Impact of Rarity and Temporality on Query Execution Quality (AVG-SCORE).

Category	GALOIS _A	GALOIS _F	GALOIS _S	GALOIS _{wo}	NL	SQL
<i>Impact of Rarity (Country)</i>						
Presidents-USA	0.570	0.652	0.587	0.089	0.352	0.320
Presidents-Venezuela	0.282	0.212	0.323	0.051	0.091	0.257
<i>Impact of Temporality</i>						
Recent	0.532	0.531	0.449	0.095	0.375	0.318
Past	0.240	0.305	0.378	0.000	0.183	0.281
All-Time	0.427	0.420	0.479	0.077	0.160	0.277

9.5 Experiment 5: Impact of SQL Query Complexity on Retrieval Quality

Objective: The objective of this experiment is to evaluate the robustness and stability of the GALOIS system as the structural and logical complexity of the input SQL query varies. Since Large Language Models (LLMs) tend to degrade in performance when handling complex logic (such as multiple joins or nested aggregations).

Methodology and Query Taxonomy: To conduct this analysis, we implemented a static classifier (`classify_query_complexity`) that parses the Abstract Syntax Tree (AST) of each SQL query via `sqlglot` and assigns it to one of six mutually exclusive complexity categories, ordered by expected difficulty:

- **SP2 (Selection-Projection Simple):** Linear queries involving only selection and projection operations with at most 2 conditions in the `WHERE` clause. These represent the simplest base case.
- **SP2> (Selection-Projection Complex):** Selection/projection queries featuring more than 2 filtering conditions, requiring the LLM to manage a denser logical context.
- **DIST (Distinct):** Queries including the `DISTINCT` clause, requiring the system to handle explicit result deduplication.
- **AGGR (Aggregation):** Queries utilizing aggregation functions (`COUNT`, `SUM`, `AVG`, `MAX`, `MIN`), which require arithmetic reasoning or counting capabilities on retrieved tuples.
- **G/O (Group By / Order By):** Queries imposing sorting or grouping of data, testing the system’s ability to structure output sequentially.

- **JOIN**: Queries involving joins between two or more tables. This is considered the most complex category, as it requires the system to decompose the problem into multiple scans (multi-table retrieval) and reconstruct relationships in local post-processing.

The GALOIS results are compared against the Ground Truth (loaded from pre-calculated JSON files) using three metrics: F1-Cell Exact, Cardinality, and Tuple Constraint of the `galois_eval.py` script. In Figure 1 we compare metrics of our run with the GALOIS paper ones and is clear that JOIN operations are worsening the results. The main reason for that is the difficulty of LLM to retrieve results that match the queries with JOIN clause, something already seen in the GALOIS results in Section 8. Instead figure 2 shows the token usage of our run compared to the GALOIS paper ones.




Figure 1: Result quality w.r.t query complexity




Figure 2: Costs w.r.t query complexity. Galois WO (not reported)

9.6 Experiment 6:

9.7 Experiment 7:

9.8 Experiment 8:

Objective: This experiment aims to quantify the efficacy of the **Logical Planner** (the component deciding which strategy to adopt) by measuring the "Optimality Gap". The Optimality Gap represents the performance difference between the automatic choice made by GALOIS_F and the theoretical optimal choice ("Oracle") that the system could have made if it had selected the best possible strategy for every single query. For every query present in the target datasets (GEO, WORLD, MOVIES), the runner executes the configurations in "brute-force" mode.

Oracle Definition and Gap Calculation: For each query q , we define the Oracle score (S_{opt}) as the maximum score obtainable among the available strategies:

$$S_{opt}(q) = \max(S_{GALOIS_A}(q), S_{GALOIS_WO}(q), S_{GALOIS_F}(q))$$

The Optimality Gap is calculated as the average difference between the Oracle and the adaptive variant:

$$Gap = \frac{1}{N} \sum_{q \in Q} (S_{opt}(q) - S_{GALOIS_F}(q))$$

A gap close to 0 indicates that the GALOIS_F automatic selector is choosing the best available strategy almost every time. A high gap suggests that, although a winning strategy exists for that query (e.g., Key-Scan), the system erroneously opted for another. Results in Figure 3 report the AVG-Score on All the queries in the dataset, and we also break the AVG-Score for each dataset.




Figure 3: AVG-Score Comparison for Galois A and Galois F against the optimal plans across all datasets.

10 Implementation Challenges & Future Work

The development of the Galois Execution Engine presented several non-trivial implementation and performance challenges, particularly concerning the **Key-Scan** operator and database resource management.

10.1 Performance Deficiencies of the Key-Scan Operator

The most significant performance challenge was encountered during the development and testing of the **Key-Scan** physical operator, where initial evaluations demonstrated **poor retrieval effectiveness** under the **GALOIS F** (Fact-checking) and **GALOIS WO** (Write-out) logical plans.

10.1.1 The Locality-of-Reference Problem

The poor performance stems from a fundamental design aspect of the **Key-Scan**'s two-phase strategy: the iterative Phase 1 is designed solely to retrieve unique primary key values ($\mathcal{K}_{all}$) without applying the initial query's non-key predicates.

For instance, consider querying the `venezuela_presidents` table with `president_name` as the primary key. In Phase 1, the LLM is prompted to retrieve *all* keys globally available in the training data (e.g., all world presidents). The actual filter condition (e.g., `country = 'Venezuela'`) is applied only in the post-retrieval Phase 2 (the Tail Request). If the LLM's initial response set $\mathcal{K}_{all}$ lacks sufficient locality (i.e., fails to yield keys related to 'Venezuela'), the final result set, post-query filtering, will be empty ($\mathcal{R}_{final} = \emptyset$). This retrieval inefficiency was common across several datasets used in the **GALOIS F** and **GALOIS WO** contexts, necessitating future work focused on improving the **Key-Scan**'s key retrieval effectiveness.

10.2 Database Resource Management for DuckDB

A critical architectural challenge involved managing the concurrent access and connection lifecycle for the **DuckDB** backend. The execution pipeline requires numerous connection instantiations across different classes for tasks such as retrieving column data types, inspecting the database schema, and executing final queries.

To address the potential for file locks and resource contention on the underlying `.duckdb` instance, the resource management for the **DuckDB** backend adheres to a strict **acquisition-release pattern**:

- **Scoped Instantiation:** Database connections are instantiated solely within the confined scope of the active schema retrieval or query execution task.
- **Guaranteed Release:** Connections are immediately released within `finally` blocks upon completion.

This explicit lifecycle management guarantees that the transition between the schema inspection phase and the subsequent in-memory virtualization phase proceeds without contention, maintaining system stability.

10.3 Addressing Sub-Optimal Prompt Structure

The core functionality of the **Table-Scan** and **Key-Scan** iterative models has been established; however, current system performance indicates that the **structure and semantic design of the prompts** generated by the existing Galois framework are sub-optimal. These prompts fail to consistently elicit high-fidelity, complete structured output from the LLMs.

This observation points to a critical area for future work: systematically improving the retrieval effectiveness by focusing on prompt engineering. Future efforts will concentrate on refining the instructions, contextualization ($\mathcal{H}$ and $\mathcal{P}_{pushdown}$ injection), and required output formats (e.g., explicit JSON schemas) to ensure the LLM provides more precise and comprehensive results. The hypothesis is that **optimized prompt semantics** will directly lead to a reduction in the number of required iterations, lower token consumption, and ultimately, higher data coverage.

11 Contributions

Alessio Demo (Badge number: 2142885)

- **First Deadline:** Developed function database connection and instance creation of DuckDB database manager for each dataset using the ingest files (.duckdb files). Developed function for retrieving the SQL queries from .sql files, processing the queries and executing them on DuckDB database instance and finally store it in a json file.
- **Second Deadline:** Implemented the Baseline NL function for interacting with the LLM with a backoff in case of gemini returns a free-plan limitation exception, save the results in FULL JSON format and next using it in the main script. Unit tests for the baseline_tools script using pytest. RAG task implementation using external libraries for creating the embeddings and evaluate the similarity between them in relation to a given prompt. Developed scripts for getting the db schema from DuckDB database for each dataset and the SQL query parsing, that will be used in the GALOIS implementation in the next tasks.
- **ThirdDeadline:** Deployed the query parser that translates SQL queries to an abstract syntax tree that makes it more easy to manipulate. Developed the Galois class that implements the logic handling the choice of physical and logical optimization techniques, also implemented the class GaloisEstimator that interacts with LLM for confidence rates involved in the process. Furthermore builded up the post-processing process that handles the execution of original query on the data frame containing the LLM response, and JOIN operations between multiple tables. Added costs optimization logic that asks the LLM to provide data only for the attributes that are involved in each query, instead of asking the entire set of attributes that make up the table. Developed Experiment 4, Experiment 5 and Experiment 8.

Writing the reports sections reporting the complete pipeline followed for developing the stuff above.

Francesco Pivotto (Badge number: 2158296)

- **First Deadline:** Implemented a modular project structure by relocating core logic to `src/` and consolidating configuration and database management logic within dedicated utility modules (`config/loaders.py`). Established the framework for multi-provider support by adding initial integration for the Grok LLM, including dedicated configuration points (`.env`, `config.yaml`, `loaders.py`).
- **Second Deadline:** Completed initial implementation of the SQL baseline, including the core logic for connecting to DuckDB and structuring query execution via the Watsonx provider. Fixed critical issues in environment variable loading, ensuring robust and standardized retrieval of secrets and settings via the central configuration object (using `loaders.py`). Modularized the `sql_baseline` module, incorporating the LCEL Pattern to prepare for dynamic LLM selection. Centralized general utility functions by moving helpers into `baseline_tools`. Introduced the dedicated `llm_wrapper` class to handle all LLM connection logic, API parameter management, and calls.
- **Third Deadline:** Implemented the core Galois Execution Engine (`GaloisExecutor`) and integrated its iterative physical operators (`Table-Scan` and `Key-Scan`).

Decoupled the entire schema management architecture into modular components (`base_schema_manager.py`, `schema_factory.py`).

Giorgia Amato (Badge number: 2159999)

- **First Deadline:** Contributed to the project initialization by setting up the repository structure, environment configuration, and shared utility scripts. Added support for `EXPLAIN` and `EXPLAIN ANALYZE` with JSON serialization, implemented automatic JSON generation for query outputs, and improved logging, debugging, and testing utilities.
- **Second Deadline:** Developed the Gemini-based SQL baseline, including IK/MC query handling, schema and raw-data injection into prompts, structured JSON parsing via Pydantic models, and robust error/timeout management. Refactored and extended unit tests for baseline tools and added new tests validating LLM interaction and Gemini-based SQL execution.
- **Third Deadline:** Implemented the iterative extensions for both SQL and NL baselines, introducing configurable execution via a `max_iter` parameter, tuple deduplication across iterations, and early-stopping conditions when no new tuples are returned. Standardized the aggregation and reporting of token usage and latency metrics for iterative executions, ensuring full compatibility with the existing JSON output schema.

Implemented the **GALOIS-WO Table-Scan** pipeline, including the `GaloisExecutor`, the iterative execution loop (initial and iterative prompts), stopping criteria, and robust JSON parsing and deduplication logic to guarantee unique tuple collection. Developed the `GaloisWOSchemaManager` for DuckDB schema inspection, table and attribute discovery, table-name resolution, key extraction, and automatic JSON schema example generation.

Designed and implemented the `demo_process_single_query` debugging utility, enabling step-by-step inspection of GALOIS execution variants (WO, A, S, F), including parser output, execution plans, intermediate debug information, statistics, and partial result inspection.

Fully implemented **Experiment 2 (Physical and Logical Optimization)**, including the experiment runner, evaluation pipeline, metric computation, result aggregation, and reproducible summary generation, ensuring strict alignment with the experimental protocol described in the GALOIS paper.

Initiated the implementation of **Experiment 6 (Evolution of Prompt Size)**, defining the experimental logic, metrics, and execution structure for measuring prompt growth across iterations in GALOIS; the completion and extension of this experiment were subsequently carried out by another team member.

Contributed to the writing and refinement of the report sections describing the iterative baselines, the GALOIS-WO Table-Scan algorithm, Experiment 2, and the corresponding numerical results.